

Deep Learning for Computer Vision

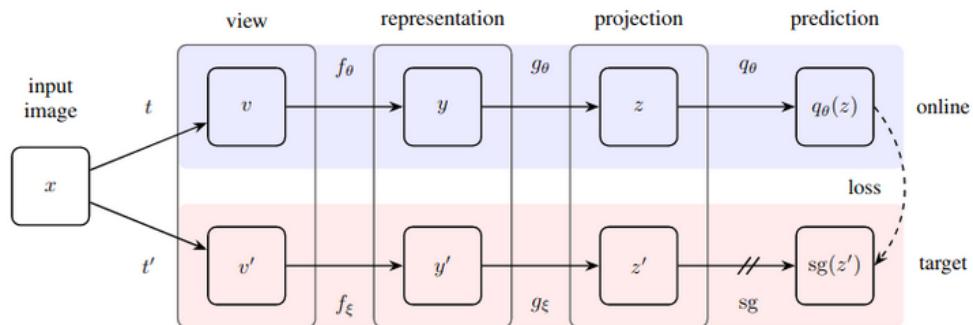
NTU, Fall 2024, Homework1

電機所碩一 李彥璋 R13921093

Problem 1: Self-Supervised Pre-training for Image Classification

1. (5%) Describe the implementation details of your SSL method for pre-training the ResNet50 backbone.

- **Name of the SSL method :** Bootstrap Your Own Latent (BYOL), 是一種無需標籤的自監督學習 (SSL) 方法，用來訓練模型。其不依賴對比學習，且無需明確的負樣本，僅通過兩個神經網絡的交互來學習有效的表徵。



- **Data augmentation :** 參考hw1_intro_2024投影片p.9, 並增加 transforms.TrivialAugmentWide()。

```
train_transform = transforms.Compose([
    transforms.Resize((128, 128)),
    transforms.CenterCrop((128, 128)),
    transforms.TrivialAugmentWide(),
    transforms.ToTensor(),
    transforms.Normalize(mean = [0.485, 0.456, 0.406], std = [0.229, 0.224, 0.225])
])
```

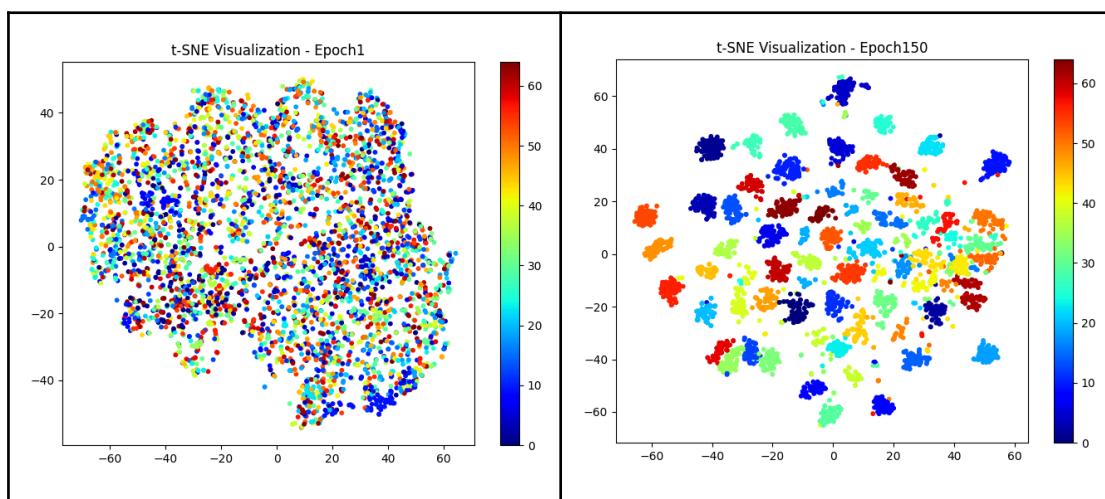
- **Learning rate schedule :** 無使用。
- **Optimizer :** learning rate 為 3e-4 的 Adam optimizer。
- **Batch size :** 64。
- **Training epoch :** 300。

2. (20%) Complete the following Table, which contains different image classification setting, and discuss/analyze the results.

Setting	Pre-training (Mini-ImageNet)	Fine-tuning (Office-Home dataset)	Validation accuracy (Office-Home dataset)
A		Train full model (backbone + classifier)	0.5246305
B	w/ label (TAs have provided this backbone)	Train full model (backbone + classifier)	0.5714285
C	w/o label (Your SSL pre-trained backbone)	Train full model (backbone + classifier)	0.5492610
D	w/ label (TAs have provided this backbone)	Fix the backbone. Train classifier only	0.3719211
E	w/o label (Your SSL pre-trained backbone)	Fix the backbone. Train classifier only	0.3694581

- 微調整個模型的效果明顯好於固定骨幹網絡(ABC v.s. DE)，這說明不僅分類器需要適應新任務，骨幹網絡也同樣需要。
- 監督學習的預訓練(B, D)效果普遍優於自監督學習的預訓練(C, E)，這可能是因為有標籤的特徵能在下游任務中表現得更加精細。
- 當完整模型經過微調之後，自監督學習的表現具有競爭力(C v.s. A)，這證明在標籤稀少的情況下，自監督學習可以成為監督預先訓練的有效替代方案。

3. (5%) Depict your visualization (t-SNE) from both the first and the last epochs. Briefly explain the results.



- 從第 1 個Epoch的零散圖形到第 150 個Epoch的明顯群集，說明模型如何隨著時間推移，在學習有意義和有組織的特徵表現方面取得進步。這代表訓練過程有效地幫助模型學習區分不同的類別或模式。

Problem 2: Semantic Segmentation

1. (3%) Draw the network architecture of your VGG16-FCN32s model (model A).

```
vgg16fcn32s(  
    (features): Sequential(  
        (0): Conv2d(3, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))  
        (1): ReLU(inplace=True)  
        (2): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))  
        (3): ReLU(inplace=True)  
        (4): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)  
        (5): Conv2d(64, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))  
        (6): ReLU(inplace=True)  
        (7): Conv2d(128, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))  
        (8): ReLU(inplace=True)  
        (9): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)  
        (10): Conv2d(128, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))  
        (11): ReLU(inplace=True)  
        (12): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))  
        (13): ReLU(inplace=True)  
        (14): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))  
        (15): ReLU(inplace=True)  
        (16): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)  
        (17): Conv2d(256, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))  
        (18): ReLU(inplace=True)  
        (19): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))  
        (20): ReLU(inplace=True)  
        (21): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))  
        (22): ReLU(inplace=True)  
        (23): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)  
        (24): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))  
        (25): ReLU(inplace=True)  
        (26): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
```

```

(27): ReLU(inplace=True)

(28): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))

(29): ReLU(inplace=True)

(30): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)

)

(fc6): Conv2d(512, 4096, kernel_size=(7, 7), stride=(1, 1), padding=(1, 1))

(relu6): ReLU(inplace=True)

(fc7): Conv2d(4096, 4096, kernel_size=(1, 1), stride=(1, 1), padding=(1, 1))

(relu7): ReLU(inplace=True)

(score_fr): Conv2d(4096, 7, kernel_size=(1, 1), stride=(1, 1), padding=(1, 1))

(relu): ReLU(inplace=True)

(upscore): ConvTranspose2d(7, 7, kernel_size=(64, 64), stride=(32, 32),
bias=False)

)

```

2. (3%) Draw the network architecture of the improved model (model B) and explain it differs from your VGG16-FCN32s model.

```

DeepLabV3(
    (backbone): IntermediateLayerGetter(
        (conv1): Conv2d(3, 64, kernel_size=(7, 7), stride=(2, 2), padding=(3, 3),
bias=False)

        (bn1): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)

        (relu): ReLU(inplace=True)

        (maxpool): MaxPool2d(kernel_size=3, stride=2, padding=1, dilation=1,
ceil_mode=False)

        (layer1): Sequential(
            (0): Bottleneck(
                (conv1): Conv2d(64, 64, kernel_size=(1, 1), stride=(1, 1), bias=False)

                (bn1): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)

                (conv2): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1),
bias=False)

```

```
(bn2): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)

(conv3): Conv2d(64, 256, kernel_size=(1, 1), stride=(1, 1), bias=False)

(bn3): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)

(relu): ReLU(inplace=True)

(downsample): Sequential(
    (0): Conv2d(64, 256, kernel_size=(1, 1), stride=(1, 1), bias=False)

    (1): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)

)
)

(1): Bottleneck(
    (conv1): Conv2d(256, 64, kernel_size=(1, 1), stride=(1, 1), bias=False)

    (bn1): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)

    (conv2): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1),
bias=False)

    (bn2): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)

    (conv3): Conv2d(64, 256, kernel_size=(1, 1), stride=(1, 1), bias=False)

    (bn3): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)

    (relu): ReLU(inplace=True)

)
)

(2): Bottleneck(
    (conv1): Conv2d(256, 64, kernel_size=(1, 1), stride=(1, 1), bias=False)

    (bn1): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)

    (conv2): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1),
bias=False)

    (bn2): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True,
```

```
track_running_stats=True)

(conv3): Conv2d(64, 256, kernel_size=(1, 1), stride=(1, 1), bias=False)

(bn3): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)

(relu): ReLU(inplace=True)

)

)

(layer2): Sequential(
(0): Bottleneck(
(conv1): Conv2d(256, 128, kernel_size=(1, 1), stride=(1, 1), bias=False)

(bn1): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)

(conv2): Conv2d(128, 128, kernel_size=(3, 3), stride=(2, 2), padding=(1, 1),
bias=False)

(bn2): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)

(conv3): Conv2d(128, 512, kernel_size=(1, 1), stride=(1, 1), bias=False)

(bn3): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)

(relu): ReLU(inplace=True)

(downsample): Sequential(
(0): Conv2d(256, 512, kernel_size=(1, 1), stride=(2, 2), bias=False)

(1): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)

)

)

(1): Bottleneck(
(conv1): Conv2d(512, 128, kernel_size=(1, 1), stride=(1, 1), bias=False)

(bn1): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)

(conv2): Conv2d(128, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1),
bias=False)
```

```
(bn2): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)

(conv3): Conv2d(128, 512, kernel_size=(1, 1), stride=(1, 1), bias=False)

(bn3): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)

(relu): ReLU(inplace=True)

)

(2): Bottleneck(
    (conv1): Conv2d(512, 128, kernel_size=(1, 1), stride=(1, 1), bias=False)

    (bn1): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)

    (conv2): Conv2d(128, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1),
bias=False)

    (bn2): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)

    (conv3): Conv2d(128, 512, kernel_size=(1, 1), stride=(1, 1), bias=False)

    (bn3): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)

    (relu): ReLU(inplace=True)

)

(3): Bottleneck(
    (conv1): Conv2d(512, 128, kernel_size=(1, 1), stride=(1, 1), bias=False)

    (bn1): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)

    (conv2): Conv2d(128, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1),
bias=False)

    (bn2): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)

    (conv3): Conv2d(128, 512, kernel_size=(1, 1), stride=(1, 1), bias=False)

    (bn3): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)

    (relu): ReLU(inplace=True)
```

```
)  
)  
(layer3): Sequential(  
(0): Bottleneck(  
(conv1): Conv2d(512, 256, kernel_size=(1, 1), stride=(1, 1), bias=False)  
(bn1): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,  
track_running_stats=True)  
(conv2): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1),  
bias=False)  
(bn2): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,  
track_running_stats=True)  
(conv3): Conv2d(256, 1024, kernel_size=(1, 1), stride=(1, 1), bias=False)  
(bn3): BatchNorm2d(1024, eps=1e-05, momentum=0.1, affine=True,  
track_running_stats=True)  
(relu): ReLU(inplace=True)  
(downsample): Sequential(  
(0): Conv2d(512, 1024, kernel_size=(1, 1), stride=(1, 1), bias=False)  
(1): BatchNorm2d(1024, eps=1e-05, momentum=0.1, affine=True,  
track_running_stats=True)  
)  
)  
(1): Bottleneck(  
(conv1): Conv2d(1024, 256, kernel_size=(1, 1), stride=(1, 1), bias=False)  
(bn1): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,  
track_running_stats=True)  
(conv2): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(2, 2),  
dilation=(2, 2), bias=False)  
(bn2): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,  
track_running_stats=True)  
(conv3): Conv2d(256, 1024, kernel_size=(1, 1), stride=(1, 1), bias=False)  
(bn3): BatchNorm2d(1024, eps=1e-05, momentum=0.1, affine=True,  
track_running_stats=True)
```

```
(relu): ReLU(inplace=True)
)
(2): Bottleneck(
    (conv1): Conv2d(1024, 256, kernel_size=(1, 1), stride=(1, 1), bias=False)
    (bn1): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,
        track_running_stats=True)
    (conv2): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(2, 2),
        dilation=(2, 2), bias=False)
    (bn2): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,
        track_running_stats=True)
    (conv3): Conv2d(256, 1024, kernel_size=(1, 1), stride=(1, 1), bias=False)
    (bn3): BatchNorm2d(1024, eps=1e-05, momentum=0.1, affine=True,
        track_running_stats=True)
    (relu): ReLU(inplace=True)
)
(3): Bottleneck(
    (conv1): Conv2d(1024, 256, kernel_size=(1, 1), stride=(1, 1), bias=False)
    (bn1): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,
        track_running_stats=True)
    (conv2): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(2, 2),
        dilation=(2, 2), bias=False)
    (bn2): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,
        track_running_stats=True)
    (conv3): Conv2d(256, 1024, kernel_size=(1, 1), stride=(1, 1), bias=False)
    (bn3): BatchNorm2d(1024, eps=1e-05, momentum=0.1, affine=True,
        track_running_stats=True)
    (relu): ReLU(inplace=True)
)
(4): Bottleneck(
    (conv1): Conv2d(1024, 256, kernel_size=(1, 1), stride=(1, 1), bias=False)
    (bn1): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,
        track_running_stats=True)
```

```
(conv2): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(2, 2),
dilation=(2, 2), bias=False)

(bn2): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)

(conv3): Conv2d(256, 1024, kernel_size=(1, 1), stride=(1, 1), bias=False)

(bn3): BatchNorm2d(1024, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)

(relu): ReLU(inplace=True)

)

(5): Bottleneck()

(conv1): Conv2d(1024, 256, kernel_size=(1, 1), stride=(1, 1), bias=False)

(bn1): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)

(conv2): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(2, 2),
dilation=(2, 2), bias=False)

(bn2): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)

(conv3): Conv2d(256, 1024, kernel_size=(1, 1), stride=(1, 1), bias=False)

(bn3): BatchNorm2d(1024, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)

(relu): ReLU(inplace=True)

)

(6): Bottleneck()

(conv1): Conv2d(1024, 256, kernel_size=(1, 1), stride=(1, 1), bias=False)

(bn1): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)

(conv2): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(2, 2),
dilation=(2, 2), bias=False)

(bn2): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)

(conv3): Conv2d(256, 1024, kernel_size=(1, 1), stride=(1, 1), bias=False)

(bn3): BatchNorm2d(1024, eps=1e-05, momentum=0.1, affine=True,
```

```
track_running_stats=True)

    (relu): ReLU(inplace=True)

)

(7): Bottleneck(
    (conv1): Conv2d(1024, 256, kernel_size=(1, 1), stride=(1, 1), bias=False)
    (bn1): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,
        track_running_stats=True)

    (conv2): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(2, 2),
        dilation=(2, 2), bias=False)

    (bn2): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,
        track_running_stats=True)

    (conv3): Conv2d(256, 1024, kernel_size=(1, 1), stride=(1, 1), bias=False)
    (bn3): BatchNorm2d(1024, eps=1e-05, momentum=0.1, affine=True,
        track_running_stats=True)

    (relu): ReLU(inplace=True)

)

(8): Bottleneck(
    (conv1): Conv2d(1024, 256, kernel_size=(1, 1), stride=(1, 1), bias=False)
    (bn1): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,
        track_running_stats=True)

    (conv2): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(2, 2),
        dilation=(2, 2), bias=False)

    (bn2): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,
        track_running_stats=True)

    (conv3): Conv2d(256, 1024, kernel_size=(1, 1), stride=(1, 1), bias=False)
    (bn3): BatchNorm2d(1024, eps=1e-05, momentum=0.1, affine=True,
        track_running_stats=True)

    (relu): ReLU(inplace=True)

)

(9): Bottleneck(
    (conv1): Conv2d(1024, 256, kernel_size=(1, 1), stride=(1, 1), bias=False)
    (bn1): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,
```

```
track_running_stats=True)

(conv2): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(2, 2),
dilation=(2, 2), bias=False)

(bn2): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)

(conv3): Conv2d(256, 1024, kernel_size=(1, 1), stride=(1, 1), bias=False)

(bn3): BatchNorm2d(1024, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)

(relu): ReLU(inplace=True)

)

(10): Bottleneck(
    (conv1): Conv2d(1024, 256, kernel_size=(1, 1), stride=(1, 1), bias=False)

    (bn1): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)

    (conv2): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(2, 2),
dilation=(2, 2), bias=False)

    (bn2): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)

    (conv3): Conv2d(256, 1024, kernel_size=(1, 1), stride=(1, 1), bias=False)

    (bn3): BatchNorm2d(1024, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)

    (relu): ReLU(inplace=True)

)

(11): Bottleneck(
    (conv1): Conv2d(1024, 256, kernel_size=(1, 1), stride=(1, 1), bias=False)

    (bn1): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)

    (conv2): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(2, 2),
dilation=(2, 2), bias=False)

    (bn2): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)

    (conv3): Conv2d(256, 1024, kernel_size=(1, 1), stride=(1, 1), bias=False)
```

```
(bn3): BatchNorm2d(1024, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)

(relu): ReLU(inplace=True)

)

(12): Bottleneck(
    (conv1): Conv2d(1024, 256, kernel_size=(1, 1), stride=(1, 1), bias=False)

    (bn1): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)

    (conv2): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(2, 2),
dilation=(2, 2), bias=False)

    (bn2): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)

    (conv3): Conv2d(256, 1024, kernel_size=(1, 1), stride=(1, 1), bias=False)

    (bn3): BatchNorm2d(1024, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)

    (relu): ReLU(inplace=True)

)

(13): Bottleneck(
    (conv1): Conv2d(1024, 256, kernel_size=(1, 1), stride=(1, 1), bias=False)

    (bn1): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)

    (conv2): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(2, 2),
dilation=(2, 2), bias=False)

    (bn2): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)

    (conv3): Conv2d(256, 1024, kernel_size=(1, 1), stride=(1, 1), bias=False)

    (bn3): BatchNorm2d(1024, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)

    (relu): ReLU(inplace=True)

)

(14): Bottleneck(
    (conv1): Conv2d(1024, 256, kernel_size=(1, 1), stride=(1, 1), bias=False)
```

```
(bn1): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)

(conv2): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(2, 2),
dilation=(2, 2), bias=False)

(bn2): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)

(conv3): Conv2d(256, 1024, kernel_size=(1, 1), stride=(1, 1), bias=False)

(bn3): BatchNorm2d(1024, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)

(relu): ReLU(inplace=True)

)

(15): Bottleneck()

(conv1): Conv2d(1024, 256, kernel_size=(1, 1), stride=(1, 1), bias=False)

(bn1): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)

(conv2): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(2, 2),
dilation=(2, 2), bias=False)

(bn2): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)

(conv3): Conv2d(256, 1024, kernel_size=(1, 1), stride=(1, 1), bias=False)

(bn3): BatchNorm2d(1024, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)

(relu): ReLU(inplace=True)

)

(16): Bottleneck()

(conv1): Conv2d(1024, 256, kernel_size=(1, 1), stride=(1, 1), bias=False)

(bn1): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)

(conv2): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(2, 2),
dilation=(2, 2), bias=False)

(bn2): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
```

```
(conv3): Conv2d(256, 1024, kernel_size=(1, 1), stride=(1, 1), bias=False)
(bn3): BatchNorm2d(1024, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
(relu): ReLU(inplace=True)
)
(17): Bottleneck(
(conv1): Conv2d(1024, 256, kernel_size=(1, 1), stride=(1, 1), bias=False)
(bn1): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
(conv2): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(2, 2),
dilation=(2, 2), bias=False)
(bn2): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
(conv3): Conv2d(256, 1024, kernel_size=(1, 1), stride=(1, 1), bias=False)
(bn3): BatchNorm2d(1024, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
(relu): ReLU(inplace=True)
)
(18): Bottleneck(
(conv1): Conv2d(1024, 256, kernel_size=(1, 1), stride=(1, 1), bias=False)
(bn1): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
(conv2): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(2, 2),
dilation=(2, 2), bias=False)
(bn2): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
(conv3): Conv2d(256, 1024, kernel_size=(1, 1), stride=(1, 1), bias=False)
(bn3): BatchNorm2d(1024, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
(relu): ReLU(inplace=True)
)
(19): Bottleneck()
```

```
(conv1): Conv2d(1024, 256, kernel_size=(1, 1), stride=(1, 1), bias=False)
(bn1): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
(conv2): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(2, 2),
dilation=(2, 2), bias=False)
(bn2): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
(conv3): Conv2d(256, 1024, kernel_size=(1, 1), stride=(1, 1), bias=False)
(bn3): BatchNorm2d(1024, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
(relu): ReLU(inplace=True)
)
(20): Bottleneck(
(conv1): Conv2d(1024, 256, kernel_size=(1, 1), stride=(1, 1), bias=False)
(bn1): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
(conv2): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(2, 2),
dilation=(2, 2), bias=False)
(bn2): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
(conv3): Conv2d(256, 1024, kernel_size=(1, 1), stride=(1, 1), bias=False)
(bn3): BatchNorm2d(1024, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
(relu): ReLU(inplace=True)
)
(21): Bottleneck(
(conv1): Conv2d(1024, 256, kernel_size=(1, 1), stride=(1, 1), bias=False)
(bn1): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
(conv2): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(2, 2),
dilation=(2, 2), bias=False)
(bn2): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,
```

```
track_running_stats=True)

(conv3): Conv2d(256, 1024, kernel_size=(1, 1), stride=(1, 1), bias=False)

(bn3): BatchNorm2d(1024, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)

(relu): ReLU(inplace=True)

)

(22): Bottleneck(
(conv1): Conv2d(1024, 256, kernel_size=(1, 1), stride=(1, 1), bias=False)

(bn1): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)

(conv2): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(2, 2),
dilation=(2, 2), bias=False)

(bn2): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)

(conv3): Conv2d(256, 1024, kernel_size=(1, 1), stride=(1, 1), bias=False)

(bn3): BatchNorm2d(1024, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)

(relu): ReLU(inplace=True)

)

)

(layer4): Sequential(
(0): Bottleneck(
(conv1): Conv2d(1024, 512, kernel_size=(1, 1), stride=(1, 1), bias=False)

(bn1): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)

(conv2): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(2, 2),
dilation=(2, 2), bias=False)

(bn2): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)

(conv3): Conv2d(512, 2048, kernel_size=(1, 1), stride=(1, 1), bias=False)

(bn3): BatchNorm2d(2048, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
```

```
(relu): ReLU(inplace=True)
(downsample): Sequential(
    (0): Conv2d(1024, 2048, kernel_size=(1, 1), stride=(1, 1), bias=False)
    (1): BatchNorm2d(2048, eps=1e-05, momentum=0.1, affine=True,
        track_running_stats=True)
)
)
(1): Bottleneck(
    (conv1): Conv2d(2048, 512, kernel_size=(1, 1), stride=(1, 1), bias=False)
    (bn1): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True,
        track_running_stats=True)
    (conv2): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(4, 4),
        dilation=(4, 4), bias=False)
    (bn2): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True,
        track_running_stats=True)
    (conv3): Conv2d(512, 2048, kernel_size=(1, 1), stride=(1, 1), bias=False)
    (bn3): BatchNorm2d(2048, eps=1e-05, momentum=0.1, affine=True,
        track_running_stats=True)
    (relu): ReLU(inplace=True)
)
(2): Bottleneck(
    (conv1): Conv2d(2048, 512, kernel_size=(1, 1), stride=(1, 1), bias=False)
    (bn1): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True,
        track_running_stats=True)
    (conv2): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(4, 4),
        dilation=(4, 4), bias=False)
    (bn2): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True,
        track_running_stats=True)
    (conv3): Conv2d(512, 2048, kernel_size=(1, 1), stride=(1, 1), bias=False)
    (bn3): BatchNorm2d(2048, eps=1e-05, momentum=0.1, affine=True,
        track_running_stats=True)
    (relu): ReLU(inplace=True)
)
```

```
)  
)  
)  
(classifier): DeepLabHead(  
    (0): ASPP(  
        (convs): ModuleList(  
            (0): Sequential(  
                (0): Conv2d(2048, 256, kernel_size=(1, 1), stride=(1, 1), bias=False)  
                (1): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,  
                    track_running_stats=True)  
                (2): ReLU()  
            )  
            (1): ASPPConv(  
                (0): Conv2d(2048, 256, kernel_size=(3, 3), stride=(1, 1), padding=(12, 12),  
                    dilation=(12, 12), bias=False)  
                (1): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,  
                    track_running_stats=True)  
                (2): ReLU()  
            )  
            (2): ASPPConv(  
                (0): Conv2d(2048, 256, kernel_size=(3, 3), stride=(1, 1), padding=(24, 24),  
                    dilation=(24, 24), bias=False)  
                (1): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,  
                    track_running_stats=True)  
                (2): ReLU()  
            )  
            (3): ASPPConv(  
                (0): Conv2d(2048, 256, kernel_size=(3, 3), stride=(1, 1), padding=(36, 36),  
                    dilation=(36, 36), bias=False)  
                (1): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,  
                    track_running_stats=True)  
                (2): ReLU()  
        )
```

```
)  
    (4): ASPPPooling(  
        (0): AdaptiveAvgPool2d(output_size=1)  
        (1): Conv2d(2048, 256, kernel_size=(1, 1), stride=(1, 1), bias=False)  
        (2): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,  
            track_running_stats=True)  
        (3): ReLU()  
    )  
    )  
    (project): Sequential(  
        (0): Conv2d(1280, 256, kernel_size=(1, 1), stride=(1, 1), bias=False)  
        (1): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,  
            track_running_stats=True)  
        (2): ReLU()  
        (3): Dropout(p=0.5, inplace=False)  
    )  
    )  
    (1): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1),  
        bias=False)  
    (2): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,  
        track_running_stats=True)  
    (3): ReLU()  
    (4): Conv2d(256, 7, kernel_size=(1, 1), stride=(1, 1))  
)  
)  
    (aux_classifier): FCNHead(  
        (0): Conv2d(1024, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1),  
            bias=False)  
        (1): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,  
            track_running_stats=True)  
        (2): ReLU()  
        (3): Dropout(p=0.1, inplace=False)  
        (4): Conv2d(256, 21, kernel_size=(1, 1), stride=(1, 1))
```

)

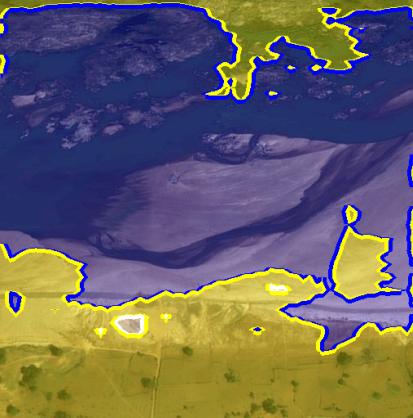
)

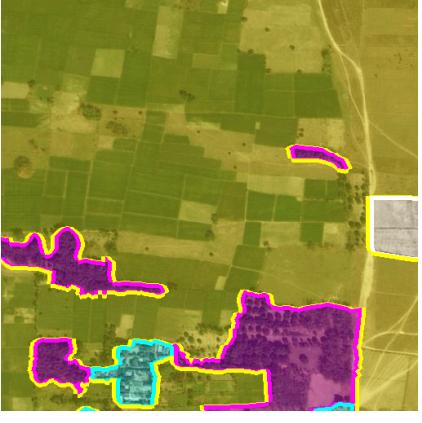
	DeepLabV3	FCN32s
Base Network Architecture	ResNet101, 101層殘差網路，使用殘差塊，深層網路訓練佳。	16層卷積網路，無使用殘差塊，深層網路訓練困難。
Semantic Segmentation Architecture	<ul style="list-style-type: none">Utilizes atrous (dilated) convolution to expand the receptive field.ASPP (Atrous Spatial Pyramid Pooling) module captures multi-scale semantic information.	<ul style="list-style-type: none">Direct upsampling, lower precision, poorer boundary handling.
Receptive Field and Contextual Information	Atrous convolution + ASPP enables capturing multi-scale contextual information.	Smaller receptive field, lacks multi-scale feature extraction.
Computational Cost	High computational cost and memory requirements.	Lightweight, lower computational cost.
Performance	Excellent at detecting boundaries and capturing contextual information, high accuracy.	Lower performance, especially in handling details compared to DeepLabV3.

3. (1%) Report mIoUs of two models on the validation set.

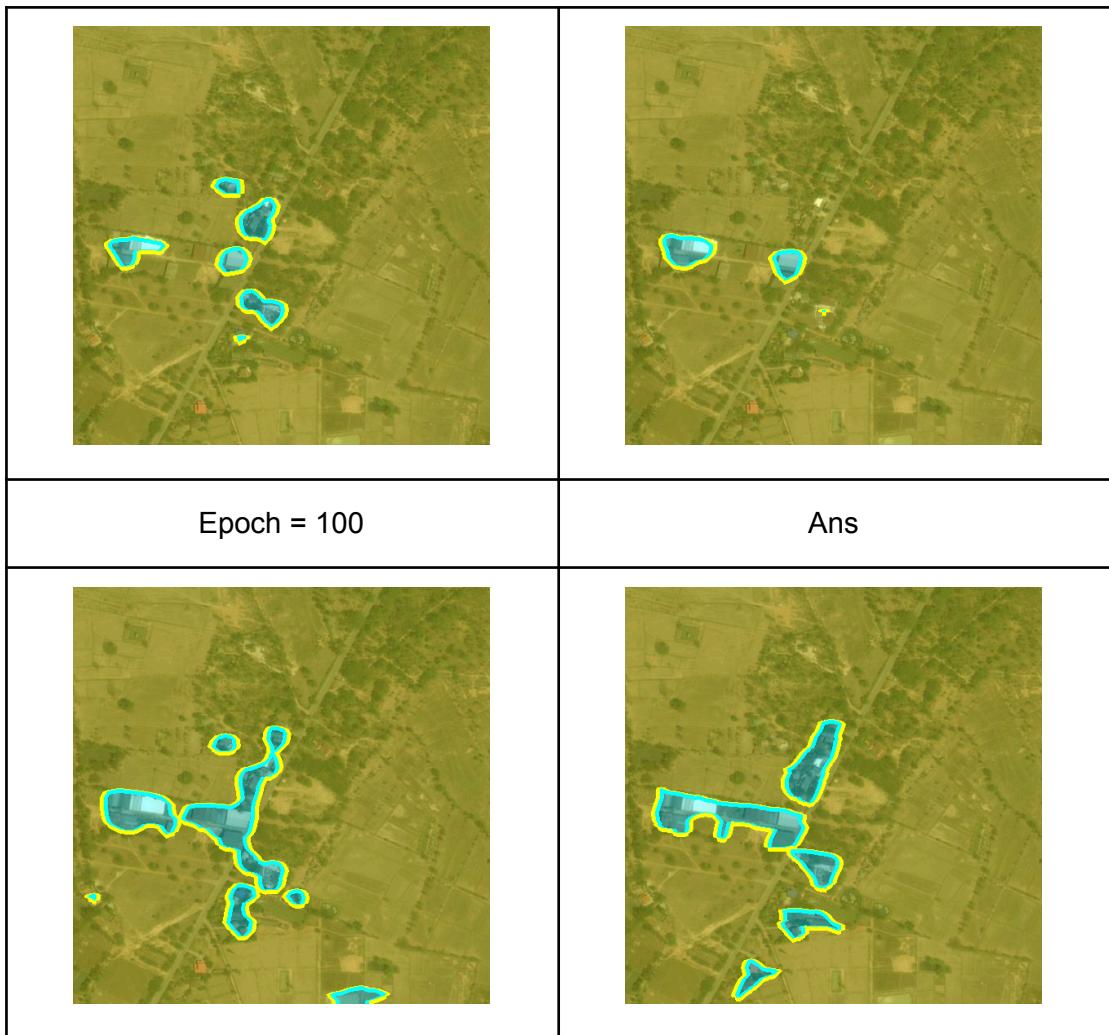
	Model A (VGG16 + FCN32s)	Model B (DeepLabV3 + ResNet101)
mIoU	0.4590993	0.7354872

4. (3%) Show the predicted segmentation mask of “validation/0013_sat.jpg”, “validation/0062_sat.jpg”, “validation/0104_sat.jpg” during the early, middle, and the final stage during the training process of the improved model.

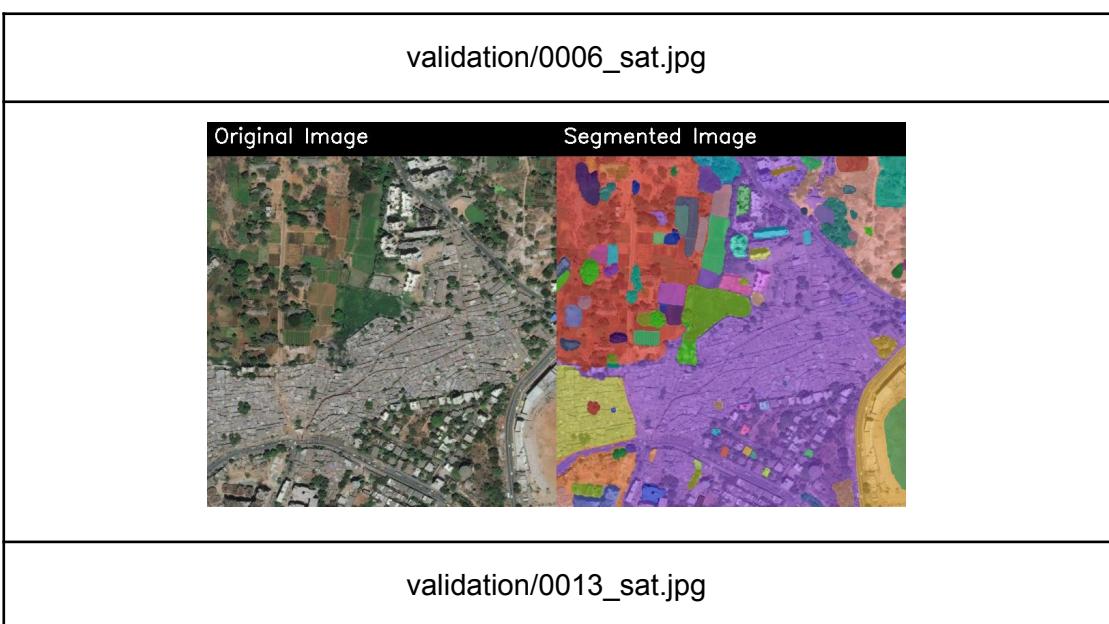
validation/ 0013 _sat.jpg	
Epoch = 1	Epoch = 50
	
Epoch = 100	Ans
	

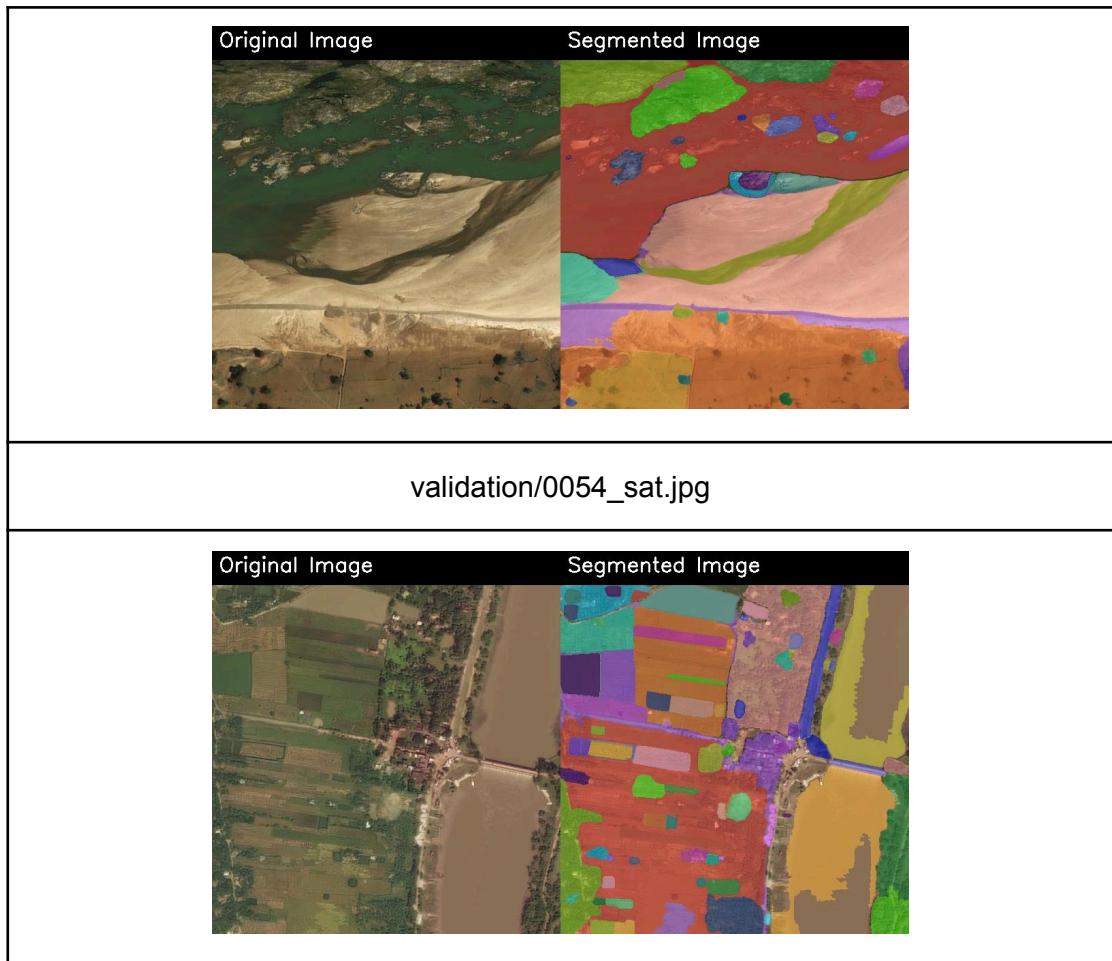
validation/0062_sat.jpg	
Epoch = 1	Epoch = 50
	
Epoch = 100	Ans
	

validation/0104_sat.jpg	
Epoch = 1	Epoch = 50



5. (10%) Use segment anything model (SAM) to segment three of the images in the validation dataset, report the result images and the method you use.





- Segment Anything Model (SAM): 使用了 Meta AI 開發的 SAM 模型來自動生成圖像的掩碼區域。這是一種針對圖像的高效分割方法，特別適合大規模的圖像處理任務。
- OpenCV (cv2): 用來讀取、處理、轉換圖像格式，並將結果儲存到硬碟上。
- Supervision Library: 用來標註分割結果和圖像，將分割後的結果展示出來。
- NumPy: 用於進行矩陣拼接和處理。

Reference

1. Pytorch官方網站
<https://pytorch.org/>
2. Bootstrap Your Own Latent (BYOL), in Pytorch <https://github.com/lucidrains/byol-pytorch>
3. Pytorch 基礎學習2_Dataset與DateLoader <https://www.wpgdadatong.com/blog/detail/46477>
4. Pytorch 基礎學習3: 圖像分類模型訓練
[https://www.wpgdadatong.com/blog/detail/46896\](https://www.wpgdadatong.com/blog/detail/46896)
5. PyTorch - 練習kaggle - Dogs vs. Cats - 使用自定義的 CNN model
<https://hackmd.io/@lido2370/S1aX6e1nN?type=view>
6. Fine-Tuning a Pre-Trained ResNet-18 Model for Image Classification with PyTorch
<https://alirezasamar.com/blog/2023/03/fine-tuning-pre-trained-resnet-18-model-image-classification-pytorch/>
7. 李宏毅 Machine Learning 2023 HW3 Image Classification
https://colab.research.google.com/drive/15A_8ilH-6-T3H0mSFrKbjDinBJI-s-16
8. Segment Anything
<https://github.com/facebookresearch/segment-anything?tab=readme-ov-file#model-checkpoints>
9. how-to-segment-anything-with-sam.ipynb
<https://github.com/roboflow/notebooks/blob/main/notebooks/how-to-segment-anything-with-sam.ipynb>
10. Chatgpt
<https://chatgpt.com/>